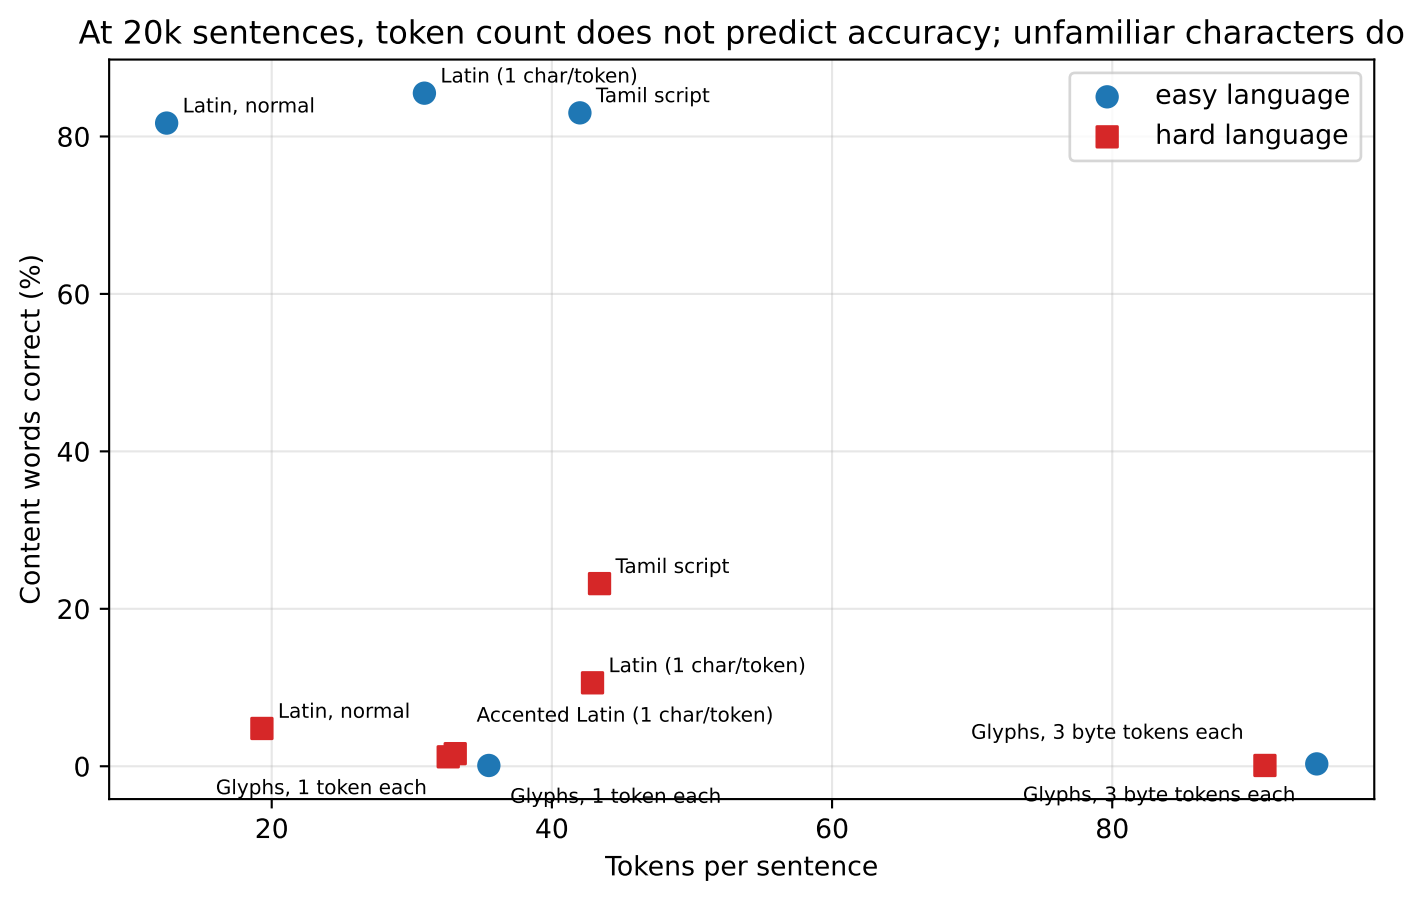
At 20k sentences, token count does not predict accuracy; unfamiliar characters do
easy language
hard language
Latin, normal
Latin (1 char/token)
Tamil script
Tamil script
Latin (1 char/token)
Latin, normal
Accented Latin (1 char/token)
Glyphs, 1 token each
Glyphs, 1 token each
Glyphs, 3 byte tokens each
Glyphs, 3 byte tokens each
Content words correct (%)
Tokens per sentence
0
20
40
60
80
20
40
60
80