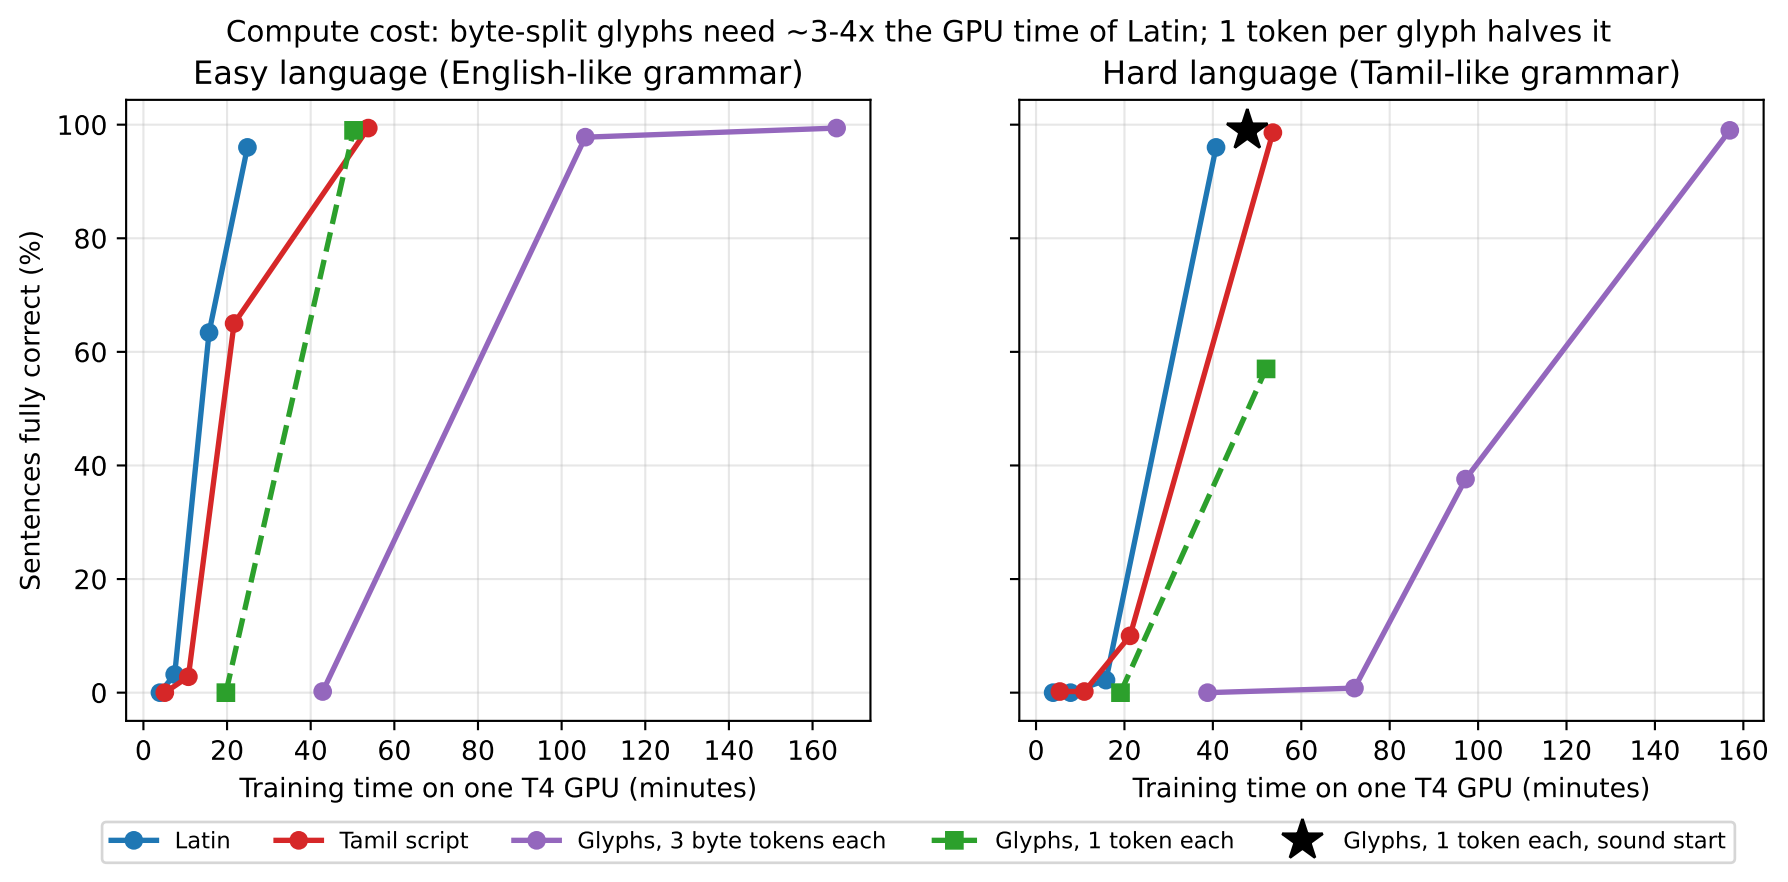
Compute cost: byte-split glyphs need ~3-4x the GPU time of Latin; 1 token per glyph halves it
Easy language (English-like grammar)
Hard language (Tamil-like grammar)
Sentences fully correct (%)
Training time on one T4 GPU (minutes)
Training time on one T4 GPU (minutes)
Latin
Tamil script
Glyphs, 3 byte tokens each
Glyphs, 1 token each
Glyphs, 1 token each, sound start